

Multiple-source shortest paths with unit weights in embedded graphs

Abstract

We describe a new algorithm that computes multiple-source shortest paths from vertices in a given boundary face to all other vertices in an embedded graph with unit weight edges.

1 Introduction

Recently, Eisanstat and Klein [?] introduced a new algorithm for computing multiple source shortest path problem for a planar graphs in linear time. Our paper attempts to generalize this idea to embedded graphs. In their paper, they maintain so-called leafmost shortest tree to get around an issue of ambiguous shortest paths between a pair of vertices. There is no direct way to generalize leafmost tree computing process to an embedded graphs, largely because there is no way to define leafmost edge of a data structure in an embedded graphs. In the following section we introduce terms that alleviate the process of ambiguous pivoting.

We can formally define multiple-source shortest paths problem as follows:

Given. Let G be a directed graph (V, \vec{E}) , embedded on a surface with genus g . All edge weights are unit.

Find. Consider boundary face f of G . $\forall v \in f$, find a shortest path to $\forall u \in V$.

Let T be the breadth first search(BFS) tree of G , and C be the BFS co-tree in G . Then there is exactly $2g$ leftover edges $L = \{e_1, e_2, \dots, e_{2g}\}$.

There exists a unique cycle λ_i in $C \cup e_i$, and $(\lambda_1, \lambda_2, \dots, \lambda_{2g}) = \Lambda$ defining homology basis. We define homological signature of an edge as follows:


$$[e]_i = \begin{cases} 1 & , \text{if } e \in \lambda_i \\ -1 & , \text{if } rev(r) \in \lambda_i \\ 0 & , \text{otherwise} \end{cases}$$

Furthermore, we define leafmost term α recursively as follows:

$$\alpha(\vec{e}^*) = \begin{cases} 1 & , \text{if } rev(e^*) \text{ is a leaf dart in } C \\ \sum_{tail(\vec{e}') = head(\vec{e}^*)} \alpha(\vec{e}') & , \text{otherwise} \end{cases}$$

We can extend above definition with $\alpha(\vec{e}) = \alpha(\vec{e}^*)$ and $\alpha(e)^* = -\alpha(rev(\vec{e}^*))$.

Let $\tilde{w}(\vec{e}) = \langle 1, [\vec{e}], \alpha(\vec{e}) \rangle$ be new weight vector for each edge in G . We refer to each component by length, homology, and leafmost terms respectively.



Definition. An edge \vec{e} is holier than \vec{e}' , if $\tilde{w}(\vec{e}) < \tilde{w}(\vec{e}')$ in lexicographic comparison. Therefore, we can define holiness of any $S \subset G$ as follows:

$$Ho(S) = \sum_{\vec{e} \in S} \tilde{w}(\vec{e})$$

2 Need of holiness

We provide a simple explanation on why it is necessary to introduce holiness to cope with ambiguity.

???

Holiest tree is a spanning tree with minimal holiness. We build Holiest tree rooted at r , using slight tweak in the Bellman-Ford algorithm for finding shortest path tree rooted at r .

BuildHoliestTree(G, \tilde{w}, r):

```

Set  $dist[r] \leftarrow \langle 0, [\vec{0}], 0 \rangle$ 
 $pred(r) \leftarrow \text{NULL}$ 
for all  $v : v \neq r$ 
   $dist[v] \leftarrow \langle \infty, [\infty], \infty \rangle$ 
   $pred(v) \leftarrow \text{NULL}$ 
put  $r$  into queue
while queue is not empty:
  Let  $u \leftarrow$  dequeue item
  for all  $u \rightarrow v$ 
    if  $v$  is not marked
      mark  $v$  and put in the queue
    if  $isTense(u \rightarrow v)$ 
      relax( $u \rightarrow v$ )

```

isTense($u \rightarrow v$):

return $dist[u] + \tilde{w}(u \rightarrow v) < dist[v]$

relax($u \rightarrow v$):

$dist[v] \leftarrow dist[u] + \tilde{w}(u \rightarrow v)$
 $pred[v] \leftarrow u$

Observation. Each vertex will be added once to the queue.

Corollary. Each edge will be relaxed at most once.

Lemma 2.1. If there is no tense edge in G , then for each $v : r \rightarrow \dots \rightarrow pred(pred(v)) \rightarrow pred(v) \rightarrow v$ is the holiest path from r to v .

Proof: Let's prove it by induction on $dist[v].length$ distance from the root r .

Base. $dist[v].length = 0$, then $v = r$, so the claim holds trivially.

Induction Step. Suppose the claim is true for all vertex $v \in V$ such that $dist[v].length < d$ for some d . Consider vertex v such that $dist[v].length = d$. By induction hypothesis, all vertices with $dist[u].length = d-1$ have holiest path correctly updated. By definition, $dist[v] = \min_{u \rightarrow v} \{dist[u] + \tilde{w}(u \rightarrow v)\}$, here $dist[u].length = d-1$. By Induction hypothesis, $dist[u]$ is not tense and can construct holiest path to u , so if there is no tense edge in G then $dist[v] = \min\{dist[u] + \tilde{w}(u \rightarrow v)\}$ holds.

□

Corollary. The algorithm will produce holiest tree rooted at r in linear time.

We now have produced our initial Holiest tree.

3 Moving Along an Edge

Consider a single edge uv , which is on the boundary face f of G . Suppose we already computed the holy-tree T_u rooted at u . We transform T_u into the holy-tree T_v as follows. First, we insert a new vertex s in the interior of the uv , bisecting it into two edges su and sv with weights:

$$w_0(s \rightarrow u) = \langle 0, [\vec{0}], 0 \rangle$$

$$w_0(s \rightarrow v) = \langle 1, [w(u \rightarrow v)], \alpha(w(u \rightarrow v)) \rangle = w(u \rightarrow v)$$

Observe that this condition implies $s = u$, therefore $T_s = T_u$. We reduce distances to u and v as follows:

$$w_\epsilon(s \rightarrow u) = \langle 0, -[w(u \rightarrow v)], -\alpha(w(u \rightarrow v)) \rangle$$

$$w_\epsilon(s \rightarrow v) = \langle 1, [\vec{0}], 0 \rangle$$

Since we reduced distance to all vertices in the graph equally, the process does not introduce any pivots. Then we define a parametric weights as follows:

$$w_\lambda(s \rightarrow u) = \langle 0, -[w(u \rightarrow v)], -\alpha(w(u \rightarrow v)) \rangle - \lambda$$

$$w_\lambda(s \rightarrow v) = \langle 1, [\vec{0}], 0 \rangle + \lambda$$

Every other dart $x \rightarrow y$ has constant parametric weight $w_\lambda(x \rightarrow y) = w(x \rightarrow y)$. We then maintain the holy tree T_λ rooted at s , with respect to the weight function w_λ , as λ increases continuously from 0 to $\langle 1, [\vec{0}], 0 \rangle$. When $\lambda = w(u \rightarrow v)$, $T_\lambda = T_v$.

In the following algorithm, **pred** defines Holy tree rooted at u , and **dist** is corresponding distance to each vertex in the graph.

MoveAlongEdge($G, u \rightarrow v, dist, pred$):

Add new vertex s
 $pred[u], pred[v] \leftarrow s$
 $\lambda \leftarrow 0$

$w(s \rightarrow u) \leftarrow \langle 0, -[w(s \rightarrow u)], -\alpha(w(s \rightarrow u)) \rangle$
 AddSubtree($\langle 0, -[w(s \rightarrow u)], -\alpha(w(s \rightarrow u)) \rangle, u$)

$w(s \rightarrow v) \leftarrow \langle 1, [\vec{0}], 0 \rangle$
 AddSubtree($\langle 0, -[w(s \rightarrow u)], -\alpha(w(s \rightarrow u)) \rangle, v$)

while $\lambda < \langle 1, [\vec{0}], 0 \rangle$:
 pivot \leftarrow FindNextPivot
 If **pivot** is non NULL **AND** $(\lambda + slack(\mathbf{pivot})/2) < \langle 1, [\vec{0}], 0 \rangle$
 Pivot(**pivot**)
 $\lambda \leftarrow \lambda + slack(\mathbf{pivot})/2$
 else
 $\delta = \langle 1, [\vec{0}], 0 \rangle - \lambda$
 AddSubtree(δ, u)
 AddSubtree($-\delta, v$)
 $\lambda \leftarrow \lambda + \delta$

4 Bounding number of pivots

We introce a clocking lemma to prove that each edge is involved in pivoting process at most $O(g)$ times.

Lemma 4.1. *Here is the claim of the clocking lemma.*

Proof: Here is the proof. □

Theorem 4.1. *Total running time of MSSP is $O(gn)$.*

Proof: Building initial holy tree takes $O(n)$ time. The process of moving around the face and pivoting takes $O(gn)$ as each dart enters and replaces $O(g)$ times. □

5 Finding pivot quickly

- What data structure do we maintain in the G^* ? Finding shortest path in network can also be understood as a Linear Programming problem as follows:
- How do we find next pivot quickly using above structure?

6 Analysis


- Building initial tree
- Pivoting
- Number of times each edge is pivoted
- Overall running time

Couple questions regarding the slack and flow:

- If the fixed tree T is arbitrary tree (not necessarily the Holiest Tree, then the flow in the answer does not have to be optimal) but solution to the linear program $\min \langle f, \text{slack}_T \rangle$ is equal to the answer from fixed tree T , not necessarily the optimal solution. That is because for each vertex, the demand satisfies the constraint and if we consider the tree T , then the value of $\min \langle f, \text{slack}_T \rangle$ would be 0, implying it is the optimal solution. (Sum cannot be negative since otherwise there is no optimal solution)
- The reason we picked the slack as the way we defined is due to the fact that slack is not negative, ensuring that the nothing bad happens.
- What makes the non-planar case special with $2g$ extra constraints?
- How does the slack in dual representation help us to find the pivots quickly?

7 Additional


NOTE: Necessity of the α definition on the edges for holiness.
Consider the following picture:



By the definition of α :

$$\alpha(p_1) - \alpha(p_2) = \sum_f \alpha(df) > 0$$

This will ensure that any two paths p_1, p_2 , whose $w(p_1) = w(p_2)$ and $[p_1]_\Lambda = [p_2]_\Lambda$, has $\alpha(p_1) \neq \alpha(p_2)$

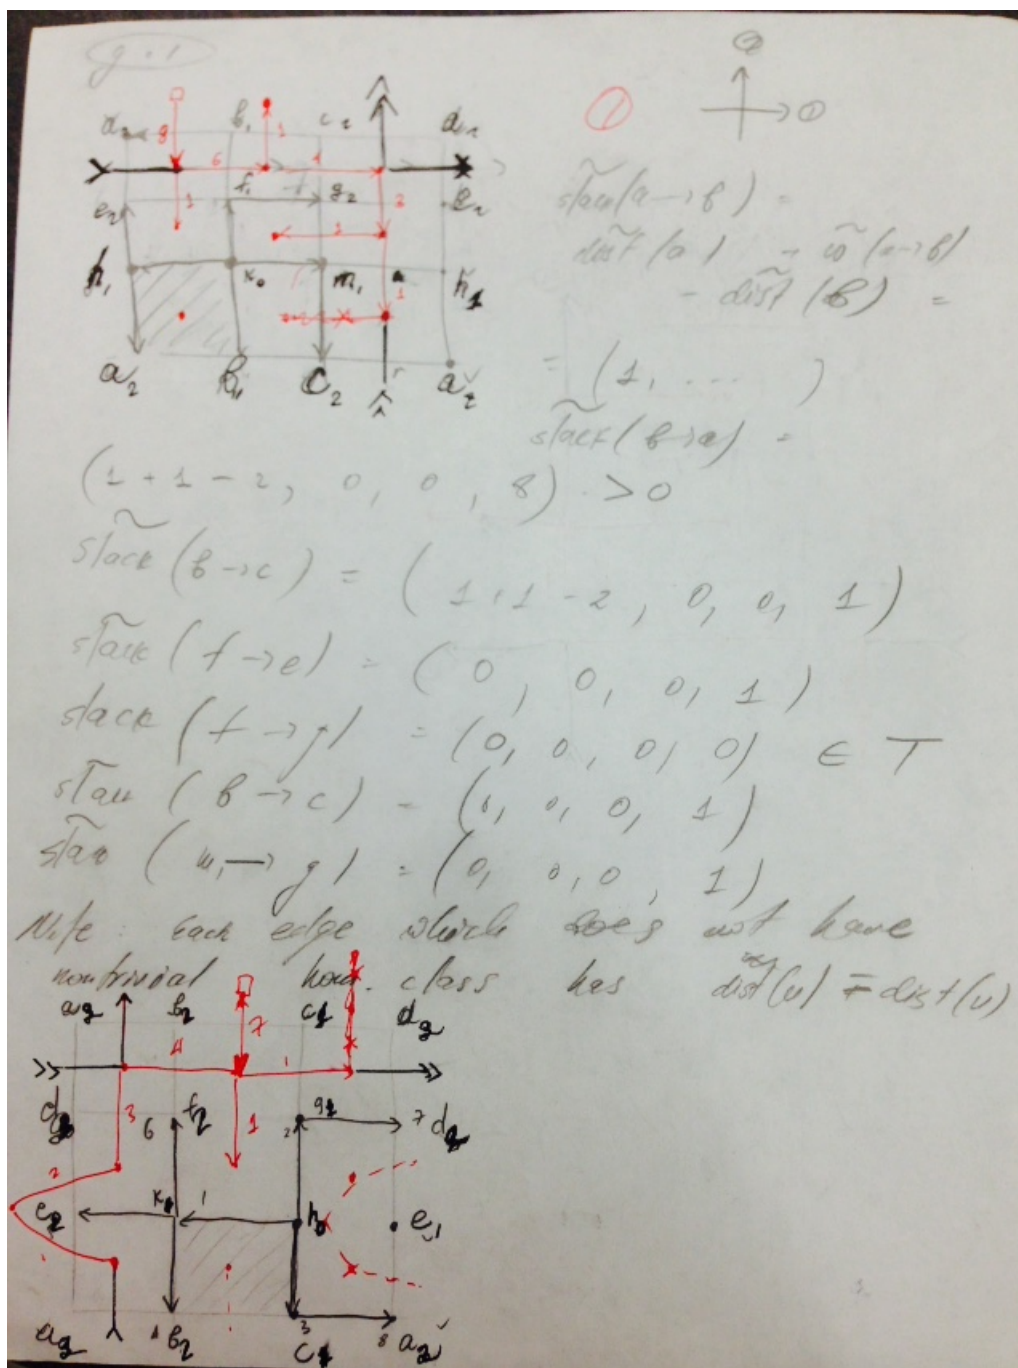


[*http : //en.wikipedia.org/wiki/Homology\(mathematics\)*](http://en.wikipedia.org/wiki/Homology(mathematics))

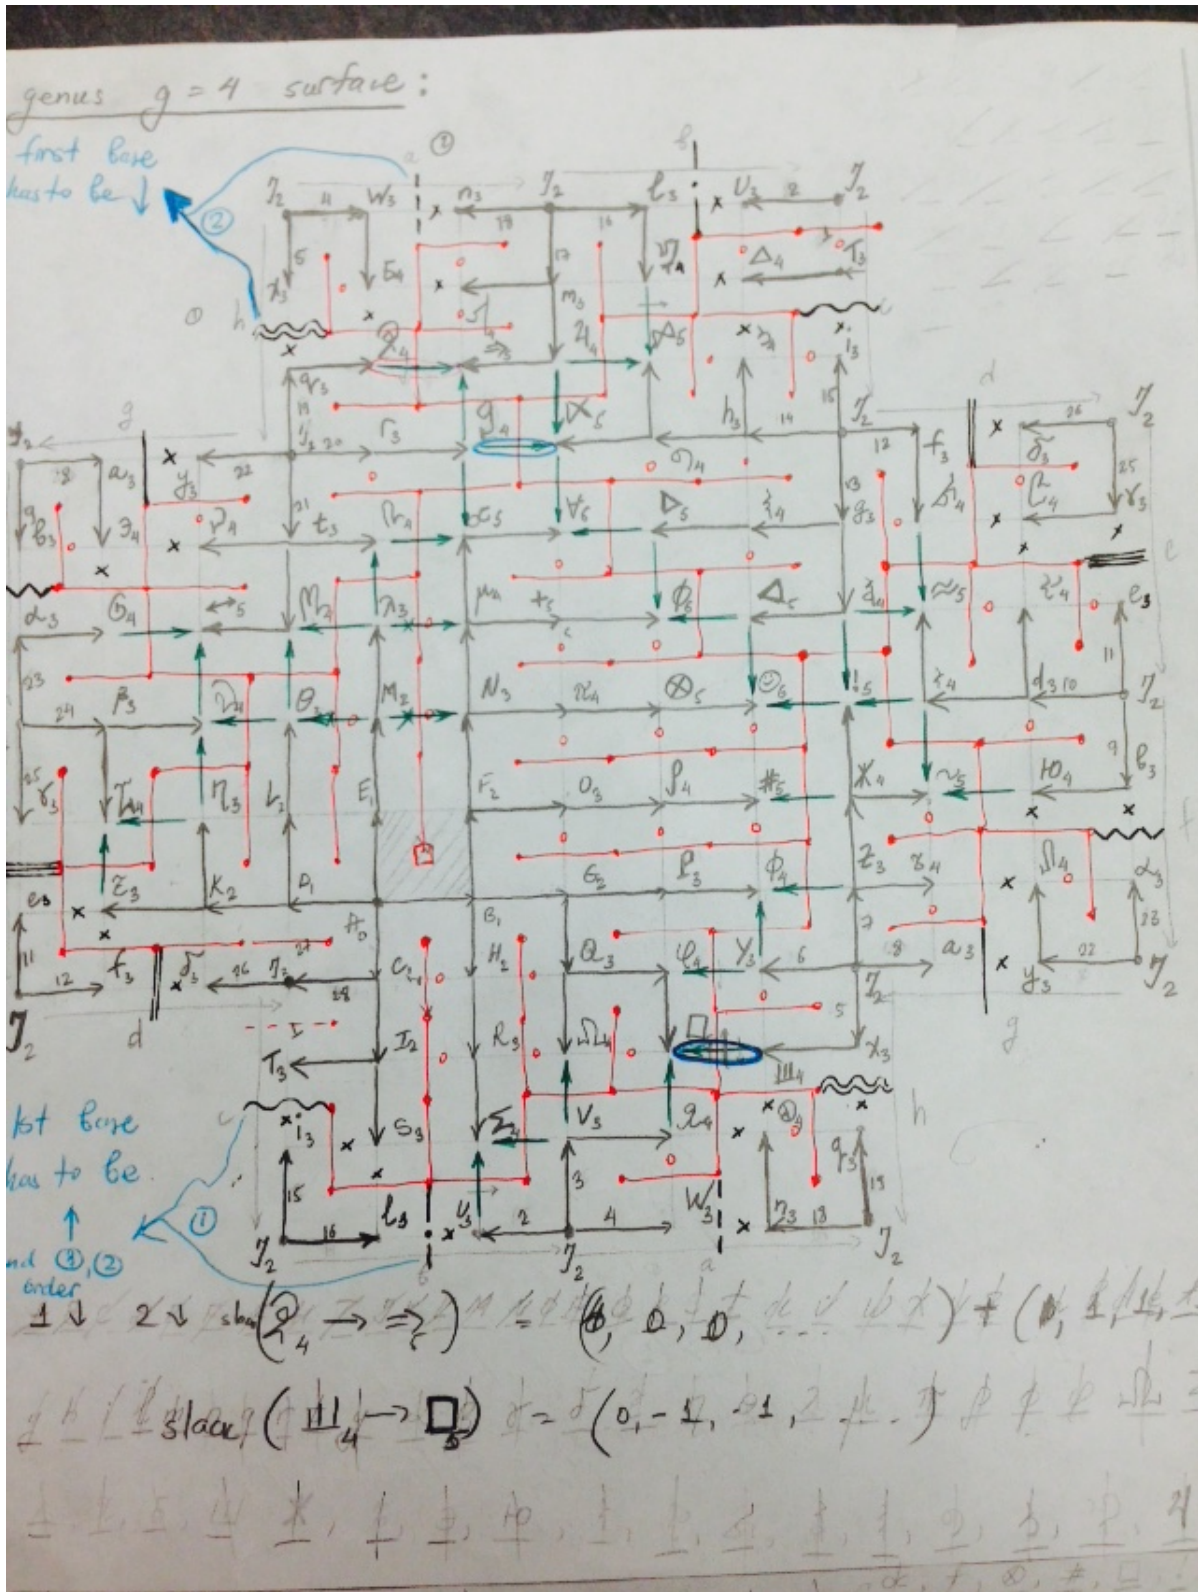
8 Working on examples:

Difference of Holiest Tree and leftmost tree:

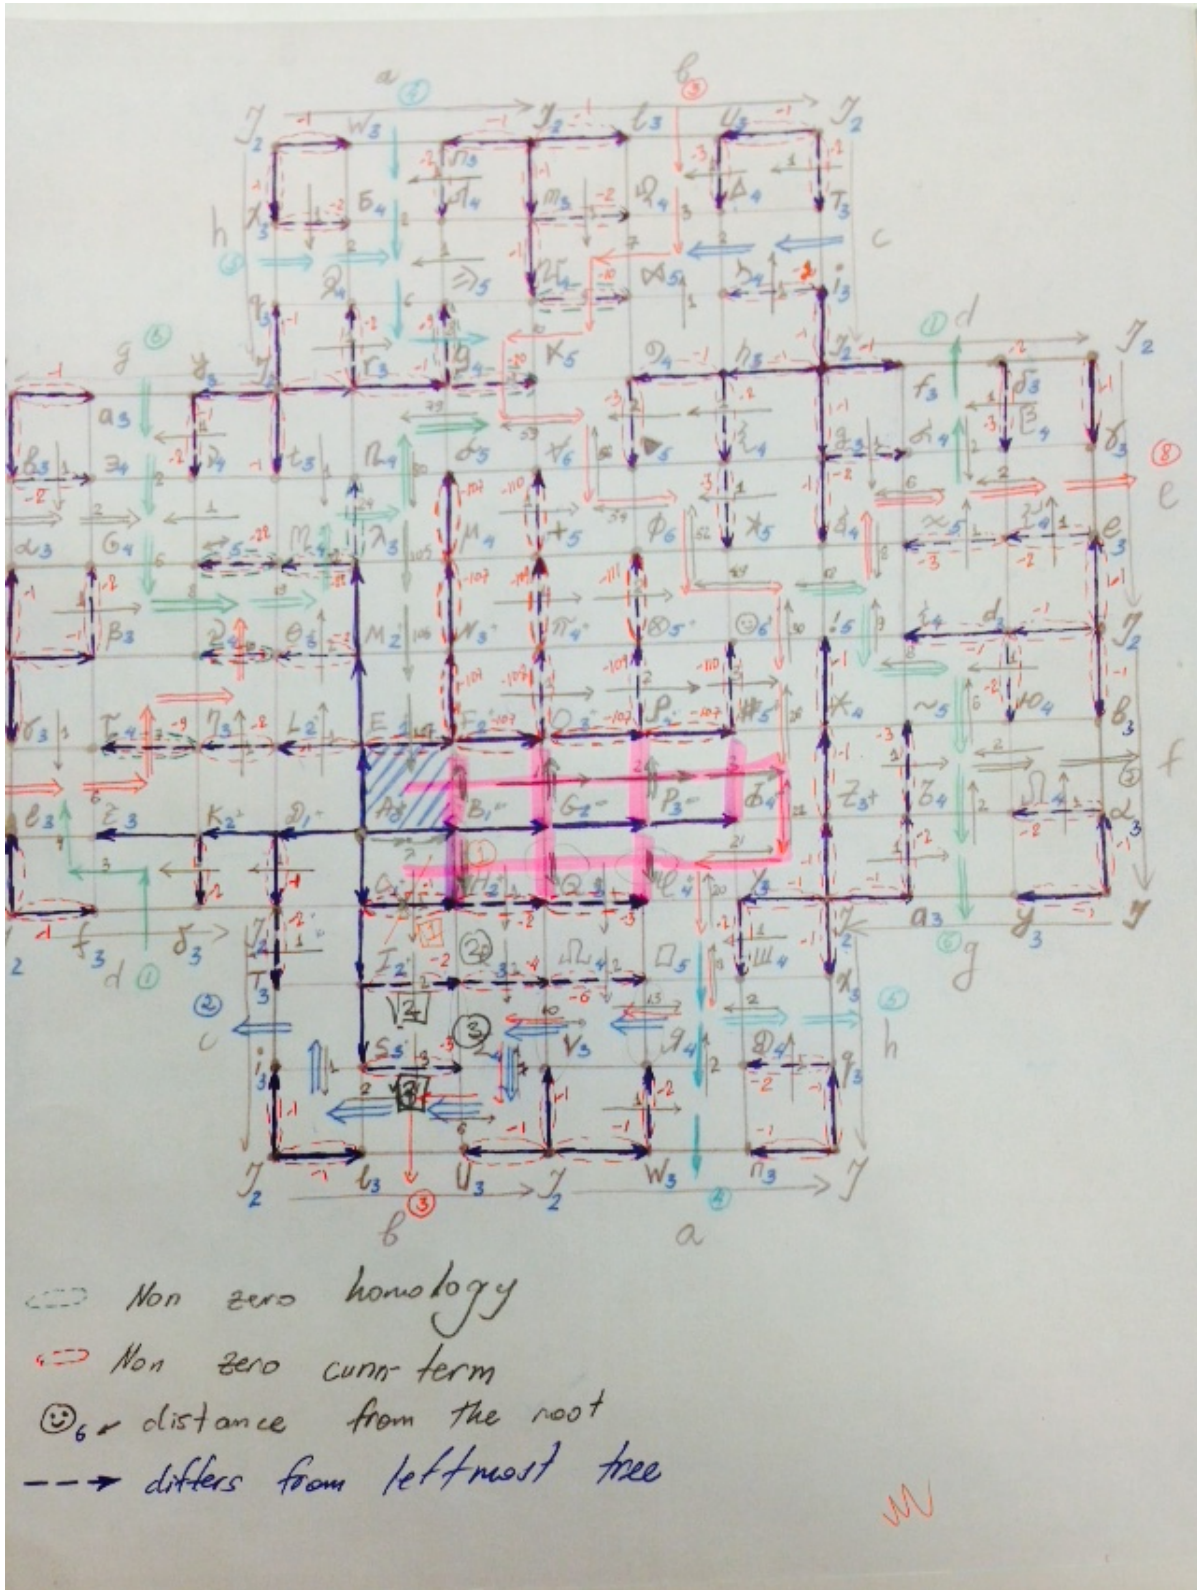
On genus $g = 1$ surface:




On genus $g = 4$ surface with initial Holy Tree build:



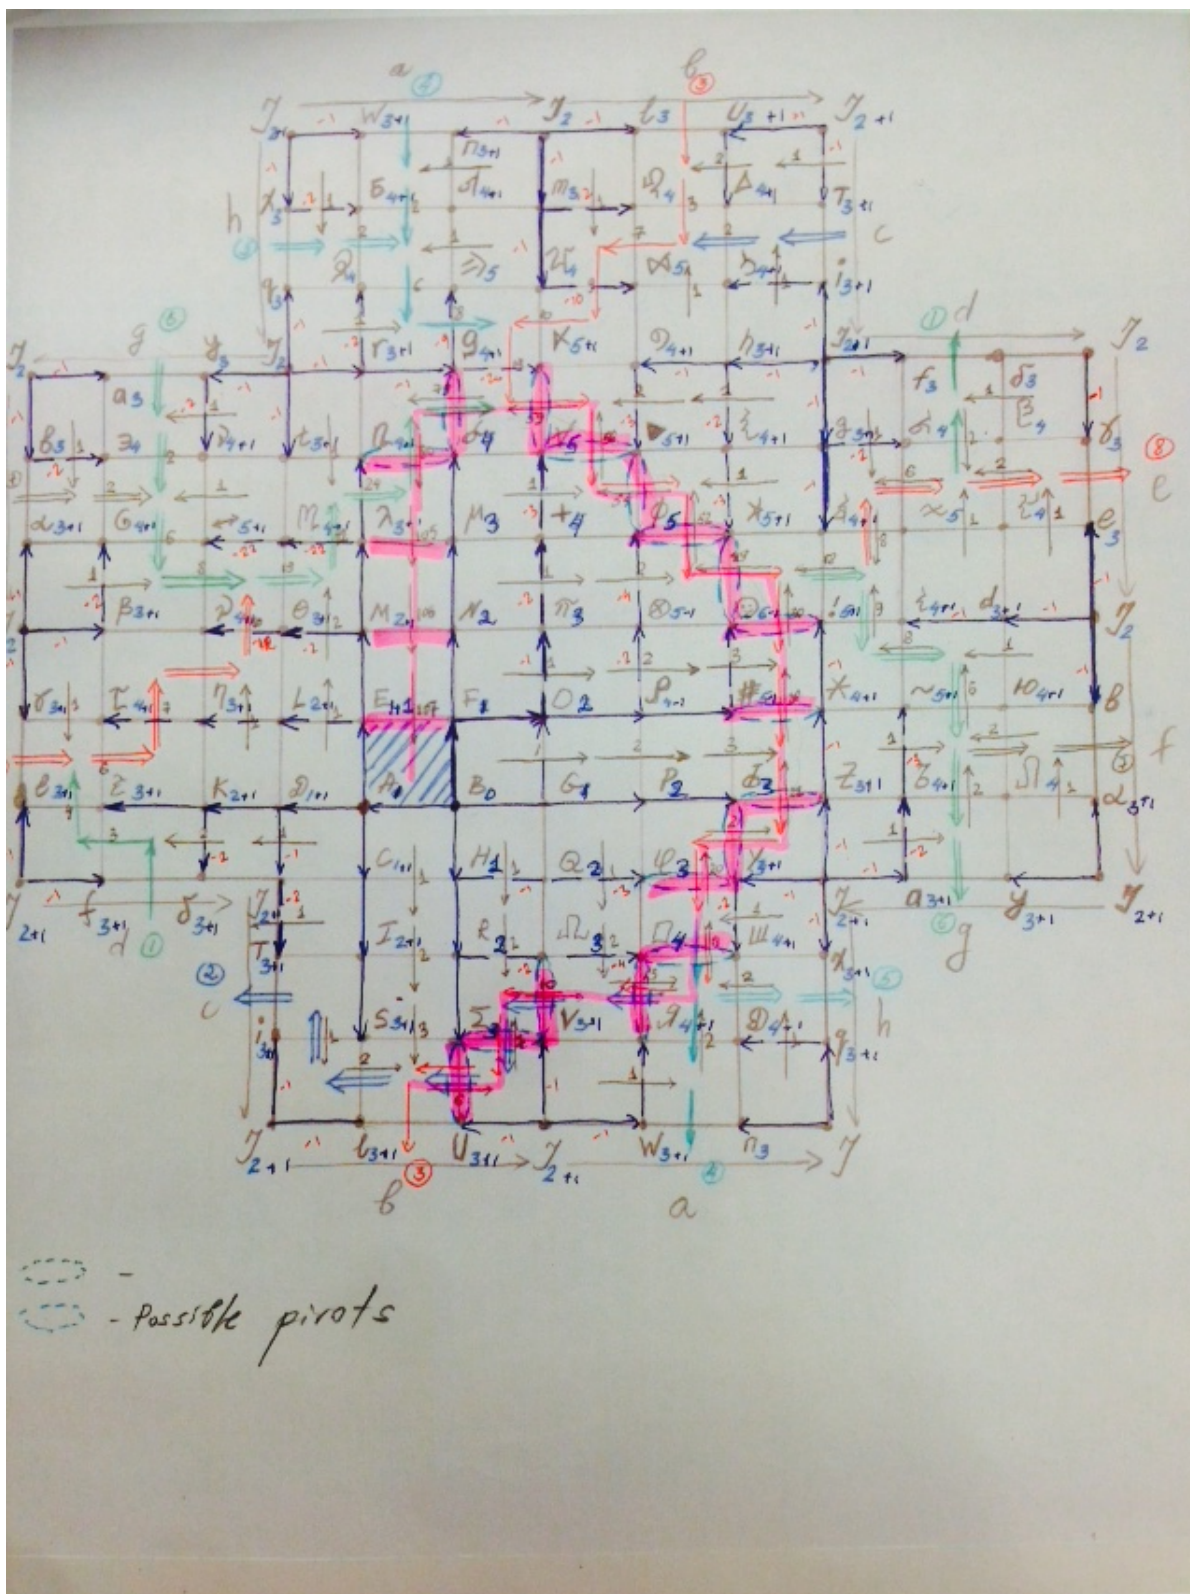
On genus $g = 4$ surface with initial pivot:



On genus $g = 4$ surface with initial pivot:



On genus $g = 5$ surface with initial pivot:



9 References:

- Cabello, Sergio, Erin W. Chambers, and Jeff Erickson. "Multiple-source shortest paths in embedded graphs." *SIAM Journal on Computing* 42.4 (2013): 1542-1571.
- Eisenstat, David, and Philip N. Klein. "Linear-time algorithms for max flow and multiple-source shortest paths in unit-weight planar graphs." *Proceedings of the forty-fifth annual ACM symposium on Theory of computing*. ACM, 2013.
- Erickson, Jeff. Maximum flows and parametric shortest paths in planar graphs. *Proceedings of the twenty-first annual ACM-SIAM symposium on Discrete Algorithms*. Society for Industrial and Applied Mathematics, 2010.